

## GPUs for AI/ML Inquiry

Eric Qian <Eric.Qian@c3.ai>

Thu 2/1/2024 5:43 PM

To:juan.perez@intelligentwaves.com <juan.perez@intelligentwaves.com>;russell.spahr@intelligentwaves.com <russell.spahr@intelligentwaves.com>;Michael Urrea <michael.urrea@intelligentwaves.com>  
Cc:Sean Summers <sean.summers@c3.ai>;Ryan McGinnity <Ryan.McGinnity@c3.ai>

Hey JP and Russ,

Regarding your inquiry of getting licenses for GPUs for AI/ML data science work, our field ops team got back to us with some steps to get that setup.

In order to create C3 worker pods to utilize the GPUs, the helm chart configurations would need to be changed. An example of that is attached to this email. The important pieces that need to be updated are the "nodeSelector" and the "nvidia.com/gpu" requests/limits in the resources section. The "nodeSelector" field in the Helm chart will depend on the GPU type (if it's not NVIDIA then it would be something else).

Depending on the type of GPU, a NVIDIA driver or something similar will likely need to be installed. Instructions here: <https://docs.nvidia.com/datacenter/cloud-native/gpu-operator/latest/getting-started.html>

In order to add new nodes to an RKE cluster, Rancher details the steps (assuming the node has a GPU already installed) here: <https://rke.docs.rancher.com/managing-clusters#:~:text=In%20order%20to%20add%20additional,list%20in%20the%20original%20cluster.>

Let us know if there are any issues or concerns and the team will be ready to assist.


Thanks,

Eric

**Eric Qian**

Associate Solution Engineer

Redwood City, California



© 2024 C3.ai. Confidential Information.